

Cliff accuracy under {visual} × {text} encoder swaps (3 seeds)

